

COL774 Assignment 4

Suchith (2021AIZ8323) and Mridul (2021AIZ8322)

2 Non-competitive part

2.1 Pre-processing

- Building on the `starter_code.py`, the captions were first enclosed within START and END tokens, and then padded with PAD tokens to make them all fixed length. Let's call this fixed length ℓ .
- The captions were tokenized at word-level. Character-level tokenization with OCR task was considered, but we didn't have the right dataset to train an OCR like model.
- After tokenization, they were converted into one-hot vectors. There were 7739 tokens.
- Images were resized to fixed size. Thresholding the images to remove the background details and converting to grayscale was tried, it didn't help.
- In images, we used the EAST text detector in opencv to crop tightly to the text. The results were poor.
- Som in conclusion, the captions were tokenized, enclosed in start and end markers and padded to fixed length ℓ after which a vocabulary mapped it to integers which was then mapped to one-hot encodings. Images were resized to fixed width-height.

2.2 CNN Encoder

We made a simple CNN block to extract features from images comprising of convolutional layers and pooling layers. Default kernel size and stride size were used.

2.3 RNN Decoder

- The output of the CNN decoder was first fed into a linear layer to create attention weights. It was then multiplied with features output by encoder.
- This was then passed into two other linear layers which would output initial hidden state and cell state to be used by RNN decoder.
- We used `torch.nn.LSTM` in bidirectional mode first, 2 layers deep. We soon realized a problem here: this only gave the final output. So, we switched to `torch.nn.LSTMCell` that would give out the hidden and cell states at each time step and it can be used to run it for as many time steps as needed.
- So, in the `torch.nn.Module.forward` method (inherited by our model) we had a `for` loop running for the length of the sequence which was the same as ℓ (see page 1).
- The hidden state at each step is seen as the unnormalized log probabilities over all possible words.
- We did teacher-forcing, just as asked for.

2.4 Loss

We used negative log likelihood loss to maximize the likelihood of each word at each time step for each example. So,

$$P_v(x^{(i)}) = [\text{softmax}(h_\theta(x^{(i)}))]_v, v \in V$$
$$LL(\theta; x, y) = \sum_{i \in M} \sum_{t=1}^T \sum_{v \in V} 1\{y_t^{(i)} = v\} \log P_v(x^{(i)})$$

where with abuse of notation, $h_\theta(x^{(i)})$ is used to mean the hidden state output by the rnn and time t by the RNN decoder. V is the set of all words, the vocabulary. M represents the mini-batch.

The parameters of the model, both encoder and decoder are then trained in an end-to-end fashion using SGD optimizer.

3 Competitive part

3.1 CNN Encoder

For the CNN Encoder, pretrained RESNET101 was used.

3.2 RNN Decoder

For the decoder, GRU based decoder was used

3.3 Loss

Negative log likelihood loss at each time step was used just as in non-competitive part with regularization.